

Inside Ezyiah

How an Australian AI bookkeeping platform turns a raw bank statement into a BAS an accountant will sign – and where the machine actually does the work.

Prepared

4 September 2026

Sources

30 public help articles · marketing site · free-tier UI

Purpose

Build reference for an independent implementation

01 *What the product actually is*

Ezyiah is a multi-tenant SaaS for Australian accounting firms and bookkeepers. A firm holds many clients; each client has bank accounts, a chart of accounts, a ledger and a set of reports. Bank data arrives either through a live Open Banking feed or as an uploaded PDF/CSV statement.


The claim it sells on is narrow and specific: the coding work that consumes the first week of every month happens without anyone doing it, and the accountant's job shifts from *producing* the books to *reviewing* them.

It is not lodgement software. It prepares, reviews and exports – the ATO submission stays with the agent.

That distinction matters commercially as well as legally: it keeps the product clear of tax-agent regulation while still charging for the expensive part of the work.

The spine: statement to report

Everything in the product hangs off one pipeline. Data enters from two sources, passes through a cascade of increasingly expensive classifiers, and forks on confidence: the machine keeps what it is sure about and hands the rest to a human. Both paths converge on the same ledger.



The order is the product. Cheap deterministic stages run before the model, so only the genuinely ambiguous tail reaches it — that is what keeps per-transaction cost viable at 600,000 transactions a year. The model *proposes*; deterministic GST and accounting validation sit downstream of it, and the confidence fork is what decides whether a human ever sees the row. The feedback edge from review back into memory is the loop the whole business model rests on.


Two properties of this diagram are worth stating plainly, because they are the difference between an accounting product and a chatbot with a spreadsheet.

The model never computes a number. It picks an account and a tax code. GST is then derived deterministically from that tax code. An LLM that calculates `gross ÷ 11` is an LLM that will eventually calculate it wrong, silently, on one transaction in ten thousand.

Nothing reaches the ledger unvalidated. Between the model and the books sit checks that the account exists, belongs to this client, is active, carries a legal tax treatment, and produces a balanced entry.

How the data is organised

Tenancy is two levels deep and everything hangs below it. A firm never sees another firm's data; within a firm, a client is the unit of work — its own workspace, own ledger, own reports.



A firm-scoped tree, with one bridge. The dashed edge is the important one: transactions do not *become* reports, they *post to* the ledger, and reports read only the ledger. That indirection is what keeps a P&L, a trial balance and a BAS agreeing with each other.

Entity type is not a dropdown


The single most consequential field in client setup is the business entity. It changes which fields the form shows, and it changes which equity accounts the client's chart of accounts is built with. Choosing wrong and discovering it after transactions are posted is an expensive mistake.

ENTITY	EXTRA FIELDS CAPTURED	EQUITY STRUCTURE GENERATED
Company	Income tax rate (default 25%), total shares, value per share	Share capital and company equity accounts
Unit Trust	Total units, value per unit	Unit-holder equity — the trust analogue of share capital
Partnership	— (partners added after save)	Per-partner share of the business
Discretionary Trust	—	Trust equity, distributions at trustee discretion
Sole Trader	—	None — the individual <i>is</i> the business

Note what is *not* asked for at creation: no email, no phone, no legal name. Five fields and you are in. Those details are collected later on the client record. That is a deliberate onboarding decision — the friction is moved out of the moment where a user might abandon.

Getting the bank data in

There are two doors, and the interesting one is the live feed. Under Australia's Consumer Data Right, the accountant cannot authorise access to their client's bank — only the client can, at their own bank. The product has to orchestrate a consent flow it does not control.



The consent never passes through the accountant. Step 6 is the one people miss: a new feed pulls *forward from today only*. Backfilling a client's history is an explicit, separate action — which is precisely the moment the "five years behind" pitch gets demonstrated.

THE MANUAL PATH, AND ITS HARD LIMIT

PDF and CSV upload, up to ten files at once. The constraint that matters: scanned or photographed PDFs do not work. The parser reads embedded text, so a statement that is really an image fails. Any competitor needs to either solve OCR or reject those files with a clear message rather than mis-parsing them silently.

ONBOARDING FRICTION WORTH KNOWING ABOUT


Business and trust accounts cannot share CDR data until someone at the business is appointed a Nominated Representative at the bank. The process differs at every institution. Ezyiah documents it bank-by-bank with support contacts — which tells you how often it blocks onboarding.

Tax codes are the hinge

Every account carries a tax code, and that code — not the account, not the amount — decides which BAS labels the transaction feeds. Get the mapping wrong and the books still balance while the BAS is quietly incorrect.

TAX CODE ON THE ACCOUNT

BAS LABEL




Direction is part of the code. A flat "GST / no GST" flag cannot produce a correct BAS, because 1A and 1B must be separated — the code has to know whether money came in or went out. The two amber rows are the ones that catch people: bank fees and interest are *input taxed*, not GST-free, and wages are *excluded* entirely and report at W1.

DISCREPANCY FOUND — WORTH AN ADVISOR'S OPINION

Ezyiah's own chart of accounts codes Interest Income as GST Free Income. Interest is a financial supply, which is *input taxed* — and the two differ on the BAS, because GST-free income reports at G1 while input-taxed sales are backed out at G4. If they are wrong, G1 is overstated for every client holding a bank account. This is exactly the class of question that needs a registered agent's sign-off rather than a developer's judgment.

The flywheel

The feature marketed as "AI Memory" is not a vector database. It is a structured rule table, written at the moment an accountant corrects something, and scoped by an explicit choice they make in the dialog.



Every correction narrows the next quarter's work. The scope choice is the mechanism: *this client only* or *every client*. Client rules outrank firm rules. Because the rule is consulted before the model runs, the flywheel cuts cost and increases consistency at the same time — the same merchant codes the same way every time, which a probabilistic classifier alone cannot guarantee.

What review actually looks like

Transactions land sorted into buckets, each showing a count of rows still awaiting attention. The accountant works a bucket until its counter clears. The interaction is deliberately industrial — this is a screen someone spends hours in.

bucket Income GST-applicable revenue	bucket Expense Business costs	bucket Other Expense GST not applicable for BAS
bucket Transfers Between the client's own accounts	bucket Subcontractors Feeds TPAR	bucket Unknown The machine declined to guess

THE FIVE ACTIONS ON A ROW

Scan similar finds every matching description so one decision can be applied in bulk. Move reassigns the account, and prompts for the memory scope. Split divides one bank line across several accounts, and will not save unless the parts sum exactly to the original. Edit adjusts date, description or amounts. Delete removes the row.


That last one is the design decision I would not copy. Deleting a posted accounting record destroys the audit trail. Void, reverse or exclude preserves history and answers the same user need.

THE METRIC THAT GOVERNS EVERYTHING

A transaction sent to review costs an accountant fifteen seconds. One confidently auto-approved into the wrong account costs them a wrong BAS. The number to optimise is therefore false auto-approval rate, not raw accuracy — and "Unknown" must always be an acceptable answer.

Ten reports, one source

Every report derives from the ledger. Three of them need an extra register first — which is why depreciation and TPAR are separate product areas rather than just another report.



One reporting layer, not ten calculators. If each report did its own arithmetic you get the classic failure: the P&L says \$50,000, the trial balance says \$49,999, and nobody can tell which is right. The registers are genuine prerequisites — a client with no subcontractors set up simply has no TPAR to generate.

The commercial model

Pricing is metered on reconciled transactions, not seats. Unlimited users, unlimited clients, unlimited bank feeds on every tier. The meter is the one thing that scales.

TIER	MONTHLY	ANNUAL	RECONCILED TRANSACTIONS / YEAR
Core	\$399	\$299	120,000
Growth	\$799	\$599	300,000
Scale	\$1,499	\$1,124	600,000
Enterprise	Custom		Multi-office, custom workflows

Unused transactions roll over. Trial is 1,000 reconciled transactions with no card. Top-up packs run from \$85 for 2,500 to \$799 for 25,000.

The strategic consequence is that gross margin is an engineering property, not a commercial one. Every transaction the deterministic layers resolve before the model is margin retained. A competitor that sends everything to an LLM has the same price list and a materially worse business.

11 *Where the gaps are*

Reading thirty help articles against the marketing site surfaces three things worth acting on.

1 · THE EXPORT PATH IS UNDOCUMENTED

The marketing site claims integration with Xero and MYOB. Across every documented workflow, "export" only ever means downloading a report. No transaction export to another ledger is described anywhere. Either it does not exist yet, or it is undocumented — and for a product positioned as a pre-processing layer *in front of* Xero, that is the load-bearing feature. It is the clearest differentiation opportunity available.

2 · DELETION OVER REVERSAL

Documented behaviour permanently deletes transactions. Production accounting systems void or reverse, because the audit trail is the product's defensibility. Choosing the safer semantics costs nothing and is genuinely better.

3 · THE DOCUMENTATION WAS SHIPPED FAST

Seven articles carry a note that sections will be added "once screenshots are available" — including the Balance Sheet's liabilities and equity, the remaining PAYG labels, and the entire With Balances view of the General Ledger. The product is moving faster than its docs, which tells you where the team's attention is.

THE ONE SENTENCE TO TAKE AWAY

The moat is not the classifier. It is the ordering — deterministic rules and learned memory ahead of the model, deterministic tax and accounting validation behind it, and a human holding the pen. That is what makes the output something a registered agent will put their name to.

Compiled from publicly available documentation and the free-tier interface for the purpose of building an independent implementation. No proprietary source, prompts, schema or

brand assets were accessed or reproduced.